

1 Conway-Maxwell Poisson (COM-Poisson, CMP) and Properties

The p.m.f. for Conway-Maxwell Poisson (CMP) is:

$$P(Y = y|\lambda, \nu) = \frac{\lambda^y}{(y!)^\nu} \frac{1}{Z(\lambda, \nu)}$$

$Z(\lambda, \nu) = Z$ is the normalization constant, i.e. $Z(\lambda, \nu) = Z = \sum_{y=0}^{\infty} \frac{\lambda^y}{(y!)^\nu}$, which doesn't have closed form in general. The domain for parameters is $\lambda, \nu > 0$ and $0 < \lambda < 1, \nu = 0$. The parameter ν controls the dispersion: 1) when $\nu = 1$, the CMP is Poisson distribution, 2) when $\nu < 1$, the distribution is over-dispersed and 3) when $\nu > 1$, the distribution is under-dispersed. When $\nu \rightarrow \infty$, the CMP approaches a Bernoulli distribution, while $\nu = 0$, it reduces to a geometric distribution.

To model the mean and dispersion simultaneously, two linear models are used for parameters λ and ν , i.e. $\log(\lambda) = \mathbf{x}'\boldsymbol{\beta}$ and $\log(\nu) = \mathbf{g}'\boldsymbol{\gamma}$, in this manuscript, I further define $\boldsymbol{\theta}' = (\boldsymbol{\beta}', \boldsymbol{\gamma}')$.

To derive the SSP for CMP regression, the keys are gradient and Hessian for log-likelihood. In the following part of this section, I will define some notations and give some necessary properties for CMP.

Assume there are n independent $Y_i \sim CMP(\lambda_i, \nu_i)$. Denote $\boldsymbol{\eta}_i = (\log(\lambda_i), \nu_i)'$, $Z_i = Z(\lambda_i, \nu_i)$ The log-likelihood for i^{th} observation is:

$$l_i(\boldsymbol{\eta}_i) = y_i \log(\lambda_i) - \log(y_i!) \nu_i - \log(Z_i)$$

Since $E(Y_i) = \frac{\partial \log(Z_i)}{\partial \log(\lambda_i)}$, $Var(Y_i) = \frac{\partial^2 \log(Z_i)}{\partial \log(\lambda_i)^2}$, $E(\log(Y_i!)) = -\frac{\partial \log(Z_i)}{\partial \nu_i}$, $Var(\log(Y_i!)) = \frac{\partial^2 \log(Z_i)}{\partial \nu_i^2}$ and $Cov(Y_i, \log(Y_i!)) = -\frac{\partial^2 \log(Z_i)}{\partial \log(\lambda_i) \partial \nu_i}$, the gradient for l_i (w.r.t $\boldsymbol{\eta}_i$) is:

$$\frac{\partial l_i}{\partial \boldsymbol{\eta}_i} = \begin{pmatrix} y_i - E(Y_i) \\ E(\log(Y_i!)) - \log(y_i!) \end{pmatrix}$$

And the Hessian for l_i (w.r.t $\boldsymbol{\eta}_i$) is:

$$\frac{\partial^2 l_i}{\partial \boldsymbol{\eta}_i \partial \boldsymbol{\eta}_i'} = \begin{pmatrix} -Var(Y_i) & Cov(Y_i, \log(Y_i!)) \\ Cov(Y_i, \log(Y_i!)) & -Var(\log(Y_i!)) \end{pmatrix}$$

The moments $E(Y_i)$, $Var(Y_i)$, $E(\log(Y_i!))$, $Var(\log(Y_i!))$ and $Cov(Y_i, \log(Y_i!))$ by using the following approximation for normalizing constant Z_i :

$$Z_i = \frac{e^{\nu_i \lambda_i^{1/\nu_i}}}{\lambda_i^{\frac{\nu_i-1}{2\nu_i}}} \left(1 + c_1(\nu_i \lambda_i^{1/\nu_i})^{-1} + c_2(\nu_i \lambda_i^{1/\nu_i})^{-2} + \mathcal{O}(\lambda_i^{\frac{-3}{\nu_i}}) \right)$$

The approximation works well when $\lambda_i \geq 2$ and $\nu_i \leq 1$, and this can be helpful when updating/ calculating the gradient and hessian matrix.

If we use models $\log(\lambda_i) = \mathbf{x}_i' \boldsymbol{\beta}$ and $\log(\nu_i) = \mathbf{g}_i' \boldsymbol{\gamma}$, by using chain rule, the gradient for l_i (w.r.t $\boldsymbol{\theta}$) is:

$$\frac{\partial l_i}{\partial \boldsymbol{\theta}} = \begin{pmatrix} [y_i - E(Y_i)] \mathbf{x}_i \\ \nu_i [E(\log(Y_i!)) - \log(y_i!)] \mathbf{g}_i \end{pmatrix}$$

And the Hessian for l_i (w.r.t $\boldsymbol{\theta}$) is:

$$\frac{\partial^2 l_i}{\partial \boldsymbol{\theta} \partial \boldsymbol{\theta}'} = \begin{pmatrix} -\text{Var}(Y_i) \mathbf{x}_i \mathbf{x}_i' & \nu_i \text{Cov}(Y_i, \log(Y_i!)) \mathbf{x}_i \mathbf{g}_i' \\ \nu_i \text{Cov}(Y_i, \log(Y_i!)) \mathbf{g}_i \mathbf{x}_i' & -\nu_i [\nu_i \text{Var}(\log(Y_i!) - E(\log(Y_i!)) + \log(y_i!))] \mathbf{g}_i \mathbf{g}_i' \end{pmatrix}$$

2 SSP for CMP

By denoting $a_i = [y_i - E(Y_i | \hat{\boldsymbol{\theta}}_{MLE})]$ and $b_i = \nu_i(\hat{\boldsymbol{\theta}}_{MLE})[E(\log(Y_i! | \hat{\boldsymbol{\theta}}_{MLE})) - \log(y_i!)]$, we define

$$\mathbf{V}_c = \frac{1}{rn^2} \sum_{i=1}^n \frac{1}{\pi_i} \begin{pmatrix} a_i^2 \mathbf{x}_i \mathbf{x}_i' & a_i b_i \mathbf{x}_i \mathbf{g}_i' \\ a_i b_i \mathbf{g}_i \mathbf{x}_i' & b_i^2 \mathbf{g}_i \mathbf{g}_i' \end{pmatrix}$$

where r is the subsample size, with subsampling probabilities π_i for all data points.

Further we denote the observed information matrix as

$$\mathbf{M}_X = \frac{1}{n} \sum_{i=1}^n \begin{pmatrix} A_i & B_i \\ B_i' & C_i \end{pmatrix}$$

where

$$\begin{aligned} A_i &= \text{Var}(Y_i | \hat{\boldsymbol{\theta}}_{MLE}) \mathbf{x}_i \mathbf{x}_i' \\ B_i &= -\nu_i(\hat{\boldsymbol{\theta}}_{MLE}) \text{Cov}(Y_i, \log(Y_i! | \hat{\boldsymbol{\theta}}_{MLE})) \mathbf{x}_i \mathbf{g}_i' \\ C_i &= \nu_i(\hat{\boldsymbol{\theta}}_{MLE}) [\nu_i(\hat{\boldsymbol{\theta}}_{MLE}) \text{Var}(\log(Y_i! | \hat{\boldsymbol{\theta}}_{MLE}) - E(\log(Y_i! | \hat{\boldsymbol{\theta}}_{MLE})) + \log(y_i!))] \mathbf{g}_i \mathbf{g}_i' \end{aligned}$$

Then follow the same steps as in OSMAC, we can show that as $n \rightarrow \infty$ and $r \rightarrow \infty$, conditional on full data matrix $\mathcal{F}_n = \mathbf{X}, \mathbf{y}$ in probability,

$$\mathbf{V}^{-1/2}(\tilde{\boldsymbol{\theta}} - \hat{\boldsymbol{\theta}}_{MLE}) \rightarrow N(0, \mathbf{I})$$

where $\mathbf{V} = \mathbf{M}_X^{-1} \mathbf{V}_c \mathbf{M}_X^{-1}$ and $\tilde{\boldsymbol{\theta}}$ is the sub-sampling estimates for $\boldsymbol{\theta}$. Then under A-optimality criterion, the sub-sampling probabilities (SSPs) π_i^{mMSE} are proportional to $\|\mathbf{M}_X^{-1} \begin{pmatrix} a_i \mathbf{x}_i \\ b_i \mathbf{g}_i \end{pmatrix}\|$, while under L-optimality criterion, $\pi_i^{mVc} \propto$

$\| \begin{pmatrix} a_i \mathbf{x}_i \\ b_i \mathbf{g}_i \end{pmatrix} \|$. In other words:

$$\pi_i^{mMSE} = \frac{\| \mathbf{M}_X^{-1} \begin{pmatrix} [y_i - E(Y_i | \hat{\boldsymbol{\theta}}_{MLE})] \mathbf{x}_i \\ \nu_i(\hat{\boldsymbol{\theta}}_{MLE}) [E(\log(Y_i!) | \hat{\boldsymbol{\theta}}_{MLE}) - \log(y_i!)] \mathbf{g}_i \end{pmatrix} \|}{\sum_{j=1}^n \| \mathbf{M}_X^{-1} \begin{pmatrix} [y_j - E(Y_j | \hat{\boldsymbol{\theta}}_{MLE})] \mathbf{x}_j \\ \nu_j(\hat{\boldsymbol{\theta}}_{MLE}) [E(\log(Y_j!) | \hat{\boldsymbol{\theta}}_{MLE}) - \log(y_j!)] \mathbf{g}_j \end{pmatrix} \|}$$

$$\pi_i^{mVc} = \frac{\| \begin{pmatrix} [y_i - E(Y_i | \hat{\boldsymbol{\theta}}_{MLE})] \mathbf{x}_i \\ \nu_i(\hat{\boldsymbol{\theta}}_{MLE}) [E(\log(Y_i!) | \hat{\boldsymbol{\theta}}_{MLE}) - \log(y_i!)] \mathbf{g}_i \end{pmatrix} \|}{\sum_{j=1}^n \| \begin{pmatrix} [y_j - E(Y_j | \hat{\boldsymbol{\theta}}_{MLE})] \mathbf{x}_j \\ \nu_j(\hat{\boldsymbol{\theta}}_{MLE}) [E(\log(Y_j!) | \hat{\boldsymbol{\theta}}_{MLE}) - \log(y_j!)] \mathbf{g}_j \end{pmatrix} \|}$$

3 Necessity for fitting CMP model


Here I first generate the data by CMP distribution.

There are $n = 10000$ independent observations, with $Y_i \sim CMP(\lambda_i, \nu_i)$ for $i = 1, 2, \dots, n$. The λ_i and ν_i are modeled as:

$$\lambda_i = \exp(x_{i1})$$

$$\nu_i = \exp(1 + g_{i1})$$

where $x_{i1} \stackrel{i.i.d.}{\sim} N(0, 1)$ and $g_{i1} \stackrel{i.i.d.}{\sim} N(1, 1)$. The following two plots show the fitted mean and median for 1) Poisson regression, 2) CMP regression, with constant $\nu_i = \nu$ (constant CMP) and 3) CMP with ν_i modeled by g_{i1} (full CMP).



The MSEs ($\frac{(Y_i - \bar{Y})^2}{n}$) for mean are 1) Poisson: 2.269, 2) constant CMP: 2.273 and 3) full CMP: 0.391. The Poisson model and constant dispersion CMP model

are not good for handling extreme observations. Therefore, it's necessary to take dispersion into account and CMP is a good choice.

The running time for these three: 1)Poisson: 0.03s, 2) constant CMP: 2.40s and 3) full CMP: 3.59s. This shows that CMP regression is much more computational expensive than regular Poisson regression, which suggests the necessity of sub-sampling.

4 Sub-sampling in CMP

To fit the CMP regression model, although we can use regular Newton-Raphson method to maximize the likelihood, the information matrix may not be stable. Therefore (as far as I know), there are two methods to deal with that: 1) use direct optimization strategy, such as L-BFGS-B and 2) optimize β and γ in an alternative way, i.e. hold one part fixed when fitting another. The alternating method reduces the problem into a two-step Newton-Raphson/ IRWLS problem.

Here I use the package implementing L-BFGS-B, and modify the objective likelihood function a bit to allow for optimization of weighted log-likelihood. To be more specific, I change:

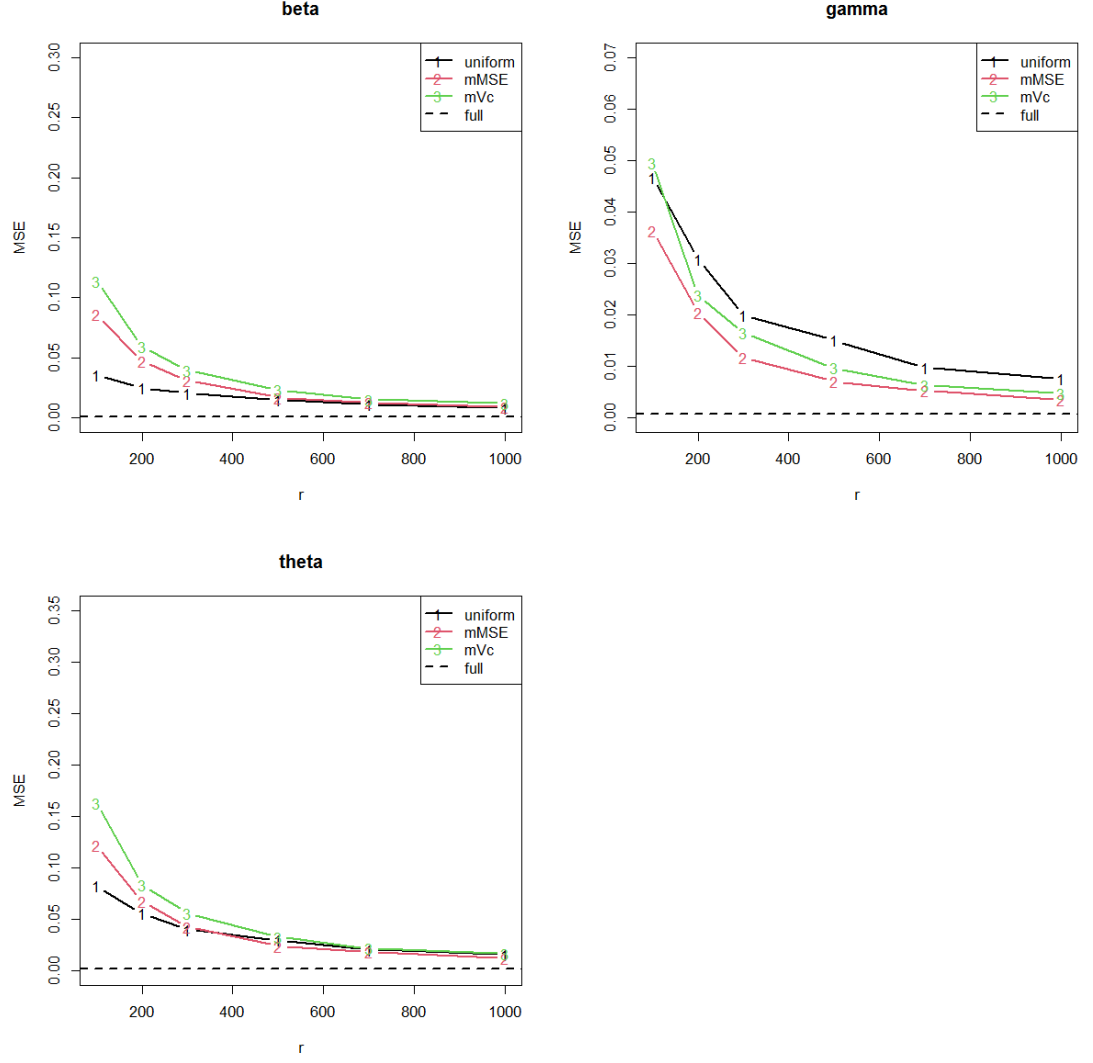
```
sum(y*log(lambda) - nu*lgamma(y+1) - logz)
```

to

```
sum(weights*(y*log(lambda) - nu*lgamma(y+1) - logz))
```

The weights are normalized such that the summation of weights equals to sample size. Therefore, I can use the modified functions to maximizes the weighted likelihood.

Then I use the bootstrap ($B = 500$) to calculate the MSE for β , γ and the overall θ . Well...



It seems that the gradient and hessian (numerically evaluated) is not stable, and the variations for evaluating gradient and hessian are larger than the variation for uniform subsampling, especially for β . (γ is always better than β , in the limited simulations)